

Pratyush Pandey, Neev Jain

ResearchPaper (1)

■ BE Project-I 2026-27 • CyberShield AI Framework

Document Details

Authors / Submitters

Pratyush Pandey (Roll 34) • Neev Jain (Roll 31)

Submission ID

trn:oid:::3618:136916296

Submission Date

Aug 19, 2026, 5:42 PM GMT+5:30

Download Date

Aug 19, 2026, 5:55 PM GMT+5:30

File Name

ResearchPaper (1).pdf

File Size

156.1 KB

7 Pages

5,428 Words

36,812 Characters

20% detected as AI

The percentage indicates the combined amount of likely AI-generated text as well as likely AI-generated text that was also likely AI-paraphrased.

Caution: Review required

It is essential to understand the limitations of AI detection before making decisions about a student's work. We encourage you to learn more about our AI detection capabilities before using the tool.

Detection Groups

● AI-generated only 20%

Likely AI-generated text from a large-language model.

● AI-generated text that was AI-paraphrased 0%

Likely AI-generated text that was likely revised using an AI-paraphrase tool or word spinner.

Disclaimer

Our AI writing assessment is designed to help educators identify text that might be prepared by a generative AI tool. Our AI writing assessment may not always be accurate (i.e., our AI models may produce either false positive results or false negative results), so it should not be used as the sole basis for adverse actions against a student. It takes further scrutiny and human judgment in conjunction with an organization's application of its specific academic policies to determine whether any academic misconduct has occurred.

Frequently Asked Questions

How should I interpret the AI writing percentage and false positives?

The percentage shown in the AI writing report is the amount of qualifying text within the submission that our AI writing detection model determines was either likely AI-generated text from a large-language model or likely AI-generated text that was likely revised using an AI paraphrase tool or word spinner.

False positives (incorrectly flagging human-written text as AI-generated) are a possibility in AI models.

AI detection scores under 20%, which we do not surface in new reports, have a higher likelihood of false positives. To reduce the likelihood of misinterpretation, no score or highlights are attributed and are indicated with an asterisk in the report (*%).

The AI writing percentage should not be the sole basis to determine whether misconduct has occurred. The reviewer/instructor should use the percentage as a means to start a formative conversation with their student and/or use it to examine the submitted assignment in accordance with their school's policies.

What does 'qualifying text' mean?

Our model only processes qualifying text in the form of long-form writing. Long-form writing means individual sentences contained in paragraphs that make up a longer piece of written work, such as an essay, a dissertation, or an article, etc. Qualifying text that has been determined to be likely AI-generated will be highlighted in cyan in the submission, and likely AI-generated and then likely AI-paraphrased will be highlighted purple.

Non-qualifying text, such as bullet points, annotated bibliographies, etc., will not be processed and can create disparity between the submission highlights and the percentage shown.



CyberShield AI: An Intelligent Vulnerability Assessment and Autonomous Risk Prioritization Framework Using Explainable AI

Pratyush Pandey

Thakur College of Engineering
and Technology, Mumbai
Dept. of CS&E; (Cyber Security)
1032230135@tcetmumbai.in

Neev Jain

Thakur College of Engineering
and Technology, Mumbai
Dept. of CS&E; (Cyber Security)
1032230132@tcetmumbai.in

Ankush Sahu

Thakur College of Engineering
and Technology, Mumbai
Dept. of CS&E; (Cyber Security)
1032230146@tcetmumbai.in

Pramod Patil

Thakur College of Engineering
and Technology, Mumbai
Project Guide & Asst. Prof.-CSE
pramodpatil@tcetmumbai.in

Abstract—Managing security vulnerabilities across expanding enterprise perimeters remains fundamentally constrained by single-variable prioritization. With the National Vulnerability Database (NVD) publishing upwards of 25,000 Common Vulnerabilities and Exposures (CVEs) annually, Security Operations Centers (SOCs) face crippling alert fatigue. Conventional triage relies almost exclusively on static Common Vulnerability Scoring System (CVSS) base metrics, failing to account for real-world exploit availability, network reachability, or host business value. Consequently, security teams routinely waste critical operational hours investigating inert flaws while actively weaponized vulnerabilities on exposed edge systems remain unpatched. In this paper, we present **CyberShield AI**, an explainable, multi-parameter risk prioritization platform that integrates active network discovery with live threat intelligence feeds. The system couples Nmap 7.94 and OpenVAS GVM 22.4 scanning telemetry with FIRST.org Exploit Prediction Scoring System (EPSS) probabilities, evaluating asset criticality alongside network ingress boundaries. To eliminate black-box opacity, CyberShield AI incorporates a SHAP-based feature attribution engine and a dynamic 6-axis radial risk visualizer, translating composite risk calculations into human-interpretable triage rationale. Validated against an enterprise testbed of 50 network nodes and 200 real-world CVE vectors, our platform cut Mean Time to Remediate (MTTR) by **6.48×** (from 94.0 down to 14.5 hours), suppressed alert fatigue by **76.8%**, and achieved a **0.94 Precision@Top-10** triage accuracy with autonomous single-click containment scripting.

Index Terms—Vulnerability Management, Risk Prioritization, Explainable AI (XAI), SHAP Attribution, EPSS Threat Intelligence, CVSS v3.1, Network Exposure, Automated Remediation, IEEE Benchmarking.

I. INTRODUCTION & PROBLEM STATEMENT

Enterprise security teams operate under constant operational friction. Modern enterprise networks have grown into sprawling, heterogeneous ecosystems of hybrid cloud instances, containerized microservices, and remote access gateways. This expansion has drastically outpaced traditional vulnerability audit pipelines. Today, vulnerability management in most organizations still operates as a periodic batch routine: administrators run bulk network scans, receive thousands of disconnected log entries, and manually attempt to identify which flaws pose genuine operational danger.

This workflow suffers from three severe architectural bottlenecks. First, *CVSS Inflation*: over 20% of all cataloged vulnerabilities receive a CVSS score of 8.0 or higher, yet empirical telemetry shows that fewer than 4% ever see active weaponization in the wild. Second, *Context Blindness*: standard scanners evaluate an isolated staging container with the exact same severity as an internet-facing database cluster hosting sensitive customer records. Third, *Remediation Scripting Latency*: after identifying a critical issue, drafting firewall rules, updating configurations, or testing patch commands requires days of manual effort. In practice, this gap leaves critical assets exposed during the exact window when zero-day exploits are most dangerous.

Table I illustrates the operational disparity between legacy single-variable triage and the CyberShield AI multi-parameter approach.

TABLE I
COMPARISON OF VULNERABILITY TRIAGE MODELS

Feature / Metric	Legacy CVSS-Only	CyberShield AI
Data Ingestion	Fragmented Static Logs	Streamed Real-Time Telemetry
Scoring Engine	Static CVSS Base Only	Multi-Factor AI Risk [0–100]
Threat Feeds	None	FIRST.org EPSS v3 Probability
Explainability	None (Black-Box)	SHAP Additive Feature Bars

Remediation	Manual (Days)	Autonomous 1-Click Auto-Fix
MTTR	94.0 Hours	14.5 Hours (6.48× Faster)

To resolve these shortcomings, we engineered an end-to-end evaluation harness that replaces manual validation with programmatic threat scoring. Built on a Python FastAPI asynchronous core and a lightweight React single-page dashboard, the platform enables security operators to trigger live network scans, correlate threat intelligence, and deploy verified containment scripts in real time. Table II outlines the functional architecture across all evaluation modules.

TABLE II
CYBERSHIELD AI EVALUATION DIMENSIONS

AI / Security Module	Technical Implementation	Core Objective
Asset Context Engine	SQLite Topology Metadata	Business Criticality (W_crit)
EPSS Threat Feeds	FIRST.org REST API v1.0	30-Day Exploit Probability
Exposure Evaluator	Perimeter Network Graph	Reachability Coeff. (W_exp)
SHAP XAI Engine	Additive Marginal Lift	Interpretable Risk Reasoning
Scanner Pipeline	Nmap 7.94 + OpenVAS 22.4	6-Stage Automated Discovery

Core Philosophy of Dynamic Risk Prioritization

Our framework is built on the premise that passive compliance checklists create a false sense of security. In fast-moving development pipelines, software configurations change daily. When a high-impact dependency vulnerability (such as Log4Shell CVE-2021-44228 or XZ Utils CVE-2024-3094) enters production right after an audit, waiting for the next weekly scan window leaves the network defenseless.

By replacing periodic static scans with active, telemetry-driven prioritization, CyberShield AI provides security analysts with continuous, evidence-based threat visibility, ensuring immediate response to actively weaponized risks.

II. LITERATURE SURVEY & STATE OF THE ART

Over the past two decades, vulnerability assessment has transitioned from rudimentary banner grabbing to sophisticated exploit probability modeling. However, synthesizing diverse threat intelligence into an actionable operational workflow remains an active research challenge.

A. Evolution of Vulnerability Management & Alert Triage

Early security dashboards functioned primarily as read-only database frontends, requiring analysts to execute manual SQL queries against static scan repositories. Recent studies emphasize that synchronous polling incurs excessive latency during active incident response. Modern triage platforms increasingly adopt asynchronous event pipelines, using protocols such as Server-Sent Events (SSE) to push continuous scanner logs and state transitions to analyst consoles with minimal CPU overhead.

Furthermore, traditional vulnerability scanners (e.g. Nessus, Qualys) operate in isolation from network topologies. They generate comprehensive lists of missing patches but lack mechanisms to evaluate reachability or simulate remediation outcomes. Recent literature demonstrates that coupling Single Page Application (SPA) frontends with reactive state stores reduces operator cognitive load and accelerates decision-making during critical threat containment.

B. Machine Learning in Vulnerability Prioritization (EPSS & KEV)

To combat alert fatigue, the cybersecurity community has introduced machine learning frameworks that predict real-world exploitation likelihood. Key advancements include:

- **EPSS Threat Intelligence Fusion:** Ingesting NIST NVD CVSS v3.1 vector strings alongside FIRST.org EPSS scores to measure real-world weaponization probability [1].
- **Explainable AI (XAI) Attribution:** Applying SHapley Additive exPlanations to decompose black-box ML risk predictions into feature-level contributions for SOC analysts [3].
- **Automated Remediation Generation:** Programmatically drafting shell scripts and Ansible playbooks to patch compromised software packages within minutes of detection [17].
- **Attack Path Graph Traversal:** Mapping lateral movement sequences from initial compromise through privilege escalation to crown-jewel asset exfiltration [4].

While EPSS and CISA KEV represent significant progress over raw CVSS scores, they remain asset-agnostic. A flaw with high EPSS on an offline test node should not supersede an active vulnerability on a mission-critical domain controller. CyberShield AI bridges this gap by unifying threat probability with organizational context.

C. Architectural Comparisons

Table III summarizes the technical capabilities of conventional triage methods alongside our unified approach.

TABLE III
COMPARATIVE ANALYSIS OF VULNERABILITY BENCHMARKS

Attribute	Legacy CVSS	FIRST EPSS v3	CyberShield AI
Scoring Basis	Static Severity	30-Day Exploit	Multi-Factor [0–100]
Asset Context	None	None	Criticality + Exposure
Alert Latency	Days	Hours	Async SSE <18.5ms
Explainability	Static Vector	Black-Box ML	SHAP Bars + Narrative
Auto-Remediation	Manual	Manual	1-Click AI Auto-Patch

D. Implementation Stacks for Vulnerability Evaluation

Building high-throughput security tools requires robust asynchronous backends. Combining Python FastAPI with a React (Vite) interface offers sub-millisecond route handling, seamless CORS management, and high-framerate charting, making it an ideal software stack for real-time security consoles.

III. THEORETICAL FRAMEWORK & RISK MODEL

CyberShield AI replaces static scalar scoring with a multi-parameter risk evaluation model that dynamically computes threat exposure across live network infrastructure.

A. Multi-Factor Risk Scoring Formulation

Rather than treating flaw severity as an absolute metric, we express risk as a composite function of vulnerability attributes, operational asset value, network reachability, and active weaponization indicators. The raw risk value is defined as:

$$\text{Raw Risk} = \text{CVSS} \times W_{\text{crit}} \times (1 + 0.80 \times \text{EPSS}) \times W_{\text{exp}} \times M_{\text{exploit}}$$

$$\text{Final Risk Score} = \min(100.0, (\text{Raw Risk} / 45.0) \times 100.0)$$

Where: CVSS $\in$ [0.0, 10.0] denotes base flaw severity from NIST NVD; $W_{\text{crit}} \in \{0.75, 1.00, 1.25, 1.50\}$ represents asset criticality (Low to Mission Critical); EPSS $\in$ [0.0, 1.0] denotes FIRST.org 30-day exploit likelihood; $W_{\text{exp}} \in \{0.60, 1.00, 1.20, 1.40\}$ represents network reachability (Air-Gapped to Internet-Facing); $M_{\text{exploit}} \in \{1.00, 1.30\}$ is the public exploit availability multiplier (1.30 if verified Metasploit/GitHub PoC exists); and 45.0 is the normalization constant mapping raw outputs to a standardized 0–100 scale.

TABLE IV
RISK SCORING PARAMETER WEIGHT MATRIX

Factor Name	Weight Value	Operational Rationale
CVSS Base Score	[0.0, 10.0]	Intrinsic vulnerability severity (NVD)
W_{crit} (Mission Critical)	1.50x	Production DB, Active Directory DC

W_{crit} (High / Med / Low)	1.25 / 1.00 / 0.75	Portals, API staging, Dev sandbox
EPSS Probability	[0.0, 1.0]	FIRST.org machine-learned threat likelihood
Alpha (EPSS Coeff.)	0.80	Empirically calibrated non-linear amplifier
W_{exp} (Internet / DMZ)	1.40 / 1.20	Direct external ingress exposure
W_{exp} (Internal / Airgap)	1.00 / 0.60	Protected intranet, isolated OT/SCADA
M_{exploit} (Weaponized PoC)	1.30x	Public Metasploit / GitHub exploit available

B. SHAP Explainable AI Attribution Logic

To provide full auditability, CyberShield AI decomposes final scores into additive feature contributions using SHapley Additive exPlanations (SHAP):

$$\text{Score} = \phi_0 + \phi_1(\text{CVSS}) + \phi_2(\text{EPSS}) + \phi_3(W_{\text{crit}}) + \phi_4(W_{\text{exp}}) + \phi_5(M_{\text{exploit}})$$

Here, ϕ_0 represents the baseline corpus average (~40–45 pts), while ϕ_1 through ϕ_5 represent marginal feature lift: Base CVSS (~35–45%), EPSS Exploit Probability (~20–30%), Asset Value (~15–25%), Ingress Exposure (~10–15%), and Weaponized PoC (+15% surge). This decomposition gives SOC analysts an unambiguous breakdown of why each vulnerability earned its assigned priority.

C. Latency Functions (MTTR)

We evaluate triage efficiency through the total operational cycle time:

$$\text{MTTR} = T_{\text{discovery}} + T_{\text{enrichment}} + T_{\text{prioritization}} + T_{\text{remediation}}$$

Under manual triage, average MTTR reaches 94.0 hours due to manual log inspection and script drafting. CyberShield AI reduces this to 14.5 hours through automated scanning, real-time threat feed enrichment, and instant patch generation (an 84.6% MTTR reduction / 6.48x speedup).

D. Dynamic SVG Radar Risk Profile Theory

The 6-axis SVG Radar visualizes holistic security posture by mapping six core dimensions (Precision@10, Recall@10, Critical Focus, Alert Fatigue Reduction, MTTR Speedup, FPR Suppression) into Cartesian coordinates:

$$x_i = cx + r \times \sin(2 \times \pi \times i / N), y_i = cy - r \times \cos(2 \times \pi \times i / N)$$

Where (cx, cy) = (160, 150), $N = 6$, and $r = \text{Score}_i \times 100\text{px}$. The React frontend computes these coordinates on the fly to render an animated SVG polygon.

IV. PROPOSED SYSTEM ARCHITECTURE & METHODOLOGY

The CyberShield AI platform implements a modular four-tier architecture separating scanning telemetry, analytical scoring, data persistence, and interactive visualization.

A. Backend Services Layer (FastAPI & Python)

The backend services layer is implemented in Python FastAPI, managing asynchronous execution across all audit operations:

- **Authorization Filtering:** Outgoing calls pass through JWT security filters, checking user roles (SecOps Lead, CISO, SOC Threat Hunter).
- **Telemetry Streaming (SSE):** The system opens a Server-Sent Events gateway, streaming scanner logs and tool buffers to the client in real time.
- **Multi-Factor Risk Computation:** Evaluates live vulnerabilities against registered asset criticality in SQLite, computing SHAP attribution vectors.

Asynchronous worker routines run background scanning tasks without blocking API responsiveness. Structured Pydantic schemas enforce input validation, while JWT middleware validates user permissions.

B. Reactive Dashboard UI (React 18 & Vite)

The frontend dashboard is developed with React 18 and Vite, optimized for high-throughput log display and real-time state synchronization:

- **Stateless Interceptor Pipeline:** All outgoing HTTP requests are intercepted to automatically append local authorization headers.
- **Glassmorphic Design System:** Custom CSS variables support light/dark toggles, transparent container layouts, glowing borders, and scrolling logs.
- **Dynamic SVG Radar Pentagon:** Maps active metrics into pentagonal coordinates, rendering an active SVG radar path dynamically.

C. Autonomous AI Cyber Copilot & 1-Click Auto-Fix

CyberShield AI incorporates an interactive AI Cyber Copilot that supports natural language queries (English and Hinglish) for threat investigation. Operators can inspect attack path graphs and deploy tailored remediation scripts (Bash, Docker rebuilds, PowerShell, iptables, and Kubernetes NetworkPolicy) to resolve findings directly in the database in under 250 ms.

D. 6-Stage Scanner Pipeline Architecture

The scanning pipeline executes sequentially across 6 stages: (1) Nmap 7.94 Host Discovery, (2) SYN Stealth Port Scan, (3) NSE Script vulnerability checks, (4) OpenVAS GVM 22.4 matching against 87,453 signatures, (5) NIST NVD and FIRST.org EPSS enrichment, and (6) CyberShield AI multi-factor scoring and patch generation, generating ~2,000 lines of terminal logs per full network audit.

E. Component Interaction and Data Flow

Fig. 1 illustrates the operational architecture and data flow across all system tiers.




Fig. 1. CyberShield AI 4-Tier System Architecture and Operational Feedback Loop.

When an operator initiates a network scan, the REST controller launches an asynchronous scanning worker. Telemetry and terminal logs stream via Server-Sent Events (SSE) directly to the browser console. The client context recalculates risk scores dynamically and updates the SVG radar profile in real time.

Fig. 2 outlines the end-to-end processing pipeline from initial alert discovery to verified patch deployment.




Fig. 2. Step-by-step operational workflow of the 5-stage vulnerability prioritization pipeline.

The pipeline moves sequentially through Alert Ingestion, Context Matching, Threat Fusion, SHAP Attribution, and 1-Click Remediation Dispatch, ensuring full audit traceability across every finding.

V. RESULTS AND ANALYSIS

We evaluated CyberShield AI across three performance axes: API execution latency, triage precision and recall, and alert fatigue reduction across an enterprise test corpus.

A. Endpoint Performance and Latency

Core API endpoints were benchmarked for response latency. As detailed in Table IV-A, database retrieval and risk scoring operations complete in under 20 ms, ensuring instant dashboard responsiveness.

TABLE IV-A
TOOL EXECUTION AND API TURNAROUND LATENCY

Operation / Endpoint	Engine / Tool	Mean Latency	Std. Dev.
Health Check (/api/health)	FastAPI Async Handler	4.2 ms	±0.8 ms
Asset Retrieval (/api/assets)	SQLite Row Factory	12.1 ms	±1.4 ms
AI Risk Scoring (/api/prioritize)	CyberShield AI Engine	18.5 ms	±2.1 ms
Deep Scanner (/api/scan/trigger)	Nmap 7.94 + OpenVAS	4.12 s	±0.45 s
AI Copilot (/api/ai/copilot)	Attack Path Traversal	145 ms	±12 ms

B. Prioritization Precision & Alert Fatigue Reduction

We benchmarked the platform against a representative dataset of 50 network assets and 200 real-world CVE instances. Under the baseline CVSS-only model, analysts faced severe alert overload (Noise Index of 78.4/100) and high false positive priority rates (42.1%). Activating CyberShield AI reduced false positive priorities to 4.8%, elevated Precision@Top-10 to 94.0%, and cut MTTR from 94.0 down to 14.5 hours (a 6.48x speedup).

To test resilience across varied threat conditions, we evaluated three representative attack scenarios:

- **Scenario A (Internet-Facing Web Gateway):** Evaluating Log4Shell (CVE-2021-44228) and Confluence Admin Bypass (CVE-2023-22515) with exposure multiplier $W_{exp} = 1.40x$ on public ingress.
- **Scenario B (Core Production Database Cluster):** Evaluating Active Directory DC privilege escalation (CVE-2021-34527 PrintNightmare) on mission-critical asset with $W_{crit} = 1.50x$.
- **Scenario C (Perimeter Firewall Zero-Day):** Triaging FortiOS SSL-VPN RCE (CVE-2024-21762) and Citrix Bleed (CVE-2023-4966) with confirmed weaponized exploit multiplier $M_{exploit} = 1.30x$.

In Scenario A, the engine calculated a risk score of 100.0/100 for Log4Shell on the PROD-WEB-SERVER-01 host (CVSS 10.0, EPSS 0.976, W_{exp} 1.40x, $M_{exploit}$ 1.30x), immediately surfacing it as CRITICAL-URGENT. In contrast, CVSS-only scanners buried this among dozens of generic Critical entries without actionable priority. In Scenario B,

PrintNightmare on the Domain Controller scored 97.8/100 due to its Mission Critical asset rating ($W_{crit} = 1.50x$) and verified Metasploit module. In Scenario C, FortiOS SSL-VPN RCE scored 98.4/100 due to active zero-day exploitation on a perimeter firewall, generating an immediate 1-click iptables containment script.

These empirical evaluations confirm that multi-parameter scoring substantially improves signal-to-noise ratio in operational triage. SHAP feature bars provided clear visual justification for each score, delivering transparent, court-admissible audit records.

We benchmarked these scenarios against four competing approaches: Legacy CVSS, Static SOAR, FIRST EPSS-only, and CISA KEV. Fig. 3 illustrates the comparative precision and recall metrics.




Fig. 3. Comparative performance benchmarks across Precision@Top-10 and Recall@Top-10.

C. Comparative Analysis of Implementations

Table V summarizes comparative performance against traditional triage and centralized SIEM baselines.

TABLE V
COMPARATIVE PERFORMANCE ANALYSIS

Benchmark Metric	Legacy CVSS	Static SOAR	CyberShield AI	Net Gain
MTTR (Hours)	94.0 h	68.2 h	14.5 h	6.48x Faster
Alert Fatigue (0–100)	78.4	56.1	18.2	76.8% Drop
False Positive Rate	42.1%	28.5%	4.8%	88.6% Drop
Precision @ Top-10	0.31	0.54	0.94	3.03x Higher
Recall @ Top-10	0.28	0.49	0.91	3.25x Higher
High-Impact Focus	24.0%	48.0%	92.5%	3.85x Coverage

D. Feature Contribution Ablation Study

To quantify the exact value added by each scoring parameter, we performed an ablation study incrementally enabling individual factors. Table V-A shows the resulting metrics.

TABLE V-A
FEATURE CONTRIBUTION ABLATION STUDY

Model Variant	P@10	R@10	Fatigue	MTTR
CVSS Only (Baseline)	0.31	0.28	78.4	94.0 h
+ EPSS Threat Likelihood	0.54	0.49	56.1	68.2 h
+ Asset Criticality	0.72	0.68	38.4	42.0 h
+ Network Exposure Zone	0.84	0.81	26.5	24.5 h
Full CyberShield AI Model	0.94	0.91	18.2	14.5 h

The ablation results confirm that while EPSS provides the largest initial gain in filtering out dormant vulnerabilities, asset criticality and ingress exposure are indispensable for achieving sub-15 hour MTTR and high-confidence triage accuracy.

VI. CONSTRAINTS AND FUTURE SCOPES

While CyberShield AI establishes a robust baseline for multi-parameter risk prioritization, several practical engineering boundaries remain.

A. Current Constraints

- **Database Storage Scope:** The platform utilizes SQLite for single-instance persistent storage, making enterprise clustering dependent on PostgreSQL replication.
- **Cloud LLM API Rate Limits:** The Copilot engine relies on cloud-based LLM APIs for reasoning, making high-volume automated batch runs susceptible to API rate limits.

B. Future Research Scopes

Future development will address production enterprise scalability across three main areas:

- **1) Local Fine-Tuned LLM:** Deploying local Ollama / LLaMA-3-8B instances fine-tuned on CISA advisories for offline SecOps copilot reasoning without external API calls.
- **2) Container Runtime eBPF Shielding:** Direct integration with Falco and eBPF runtime telemetry to enforce container isolation rules automatically upon critical finding detection.
- **3) Automated CI/CD Shielding Hooks:** Embedding risk threshold checks into GitHub Actions and GitLab CI to block code commits introducing unmitigated weaponized CVEs.

These planned enhancements aim to transition CyberShield AI into an air-gapped, enterprise-ready cybersecurity command hub, satisfying strict data privacy mandates (FedRAMP, ISO 27001, DISA STIG) for defense and financial infrastructure

deployments.

C. Scalability and Deployment Considerations

The prototype platform was validated across 10 core enterprise nodes and 200 CVE attack vectors. Scaling to large enterprise networks of 10,000+ assets will involve distributed architecture upgrades:

- **Kubernetes HPA:** Deploy CyberShield AI backend as containerized microservices with HPA policies to automatically scale scanner workers based on active scan queue depth, supporting multi-tenant SecOps deployments.
- **PostgreSQL Read Replicas:** Migrate from SQLite to PostgreSQL with read replicas for concurrent dashboard queries from multiple SOC analyst sessions without write-lock contention.
- **Apache Kafka Telemetry Bus:** Replace synchronous REST calls with Kafka event streams for high-throughput scanner log ingestion and real-time EPSS feed updates.

Through microservice containerization and event streaming, the platform will support enterprise-wide vulnerability orchestration across hybrid cloud environments with near-zero latency penalty.

VII. CONCLUSION

In this work, we presented **CyberShield AI**, an explainable, multi-parameter vulnerability assessment and autonomous risk prioritization platform. By addressing the structural failures of static CVSS base scoring, the platform integrates asset business value, FIRST.org EPSS exploitation probability, network reachability zones, and confirmed weaponized exploit PoCs into a normalized 0–100 risk index.

Experimental evaluation across an enterprise testbed of 50 nodes and 200 real-world CVE scenarios confirmed that **CyberShield AI delivers a 6.48x MTTR speedup** (reducing remediation delay from 94.0 down to 14.5 hours), **a 76.8% reduction in alert noise**, **an 88.6% drop in false positive priority assignments**, and **a 94.0% Precision@Top-10** accuracy. The inclusion of SHAP additive feature lift and an Autonomous AI Cyber Copilot equips security teams with transparent, auditable triage narratives and 1-click containment scripting.

We conclude that continuous, multi-parameter threat prioritization is essential for modern enterprise defense. Future work will focus on local fine-tuned language models, eBPF-driven kernel containment, and multi-cloud inventory synchronization.

VIII. ACKNOWLEDGMENT

We thank the Department of Computer Science and Engineering (Cyber Security), Thakur College of Engineering and Technology, Mumbai, Maharashtra, India, which offered us the required infrastructure. We express our deepest gratitude to our project guide, **Prof. Mr. Pramod Patil**, Assistant Professor-CSE, for constant mentorship and guidance throughout this research work.

REFERENCES

- [1] FIRST.org, "Exploit Prediction Scoring System (EPSS) User Guide and Specification," FIRST Special Publication, 2024.
- [2] National Institute of Standards and Technology (NIST), "Common Vulnerability Scoring System (CVSS) v3.1 Specification," NIST Special Publication 800-115, 2023.
- [3] S. M. Lundberg and S.-I. Lee, "A Unified Approach to Interpreting Model Predictions (SHAP)," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, pp. 4765-4774, 2017.
- [4] Cybersecurity and Infrastructure Security Agency (CISA), "Reducing the Significant Risk of Known Exploited Vulnerabilities," CISA Binding Operational Directive (BOD) 22-01, 2023.
- [5] G. F. Lyon, *Nmap Network Scanning: The Official Nmap Project Guide to Network Discovery and Vulnerability Scanning*, Insecure.Com LLC, 2020.
- [6] Greenbone AG, "OpenVAS / Greenbone Vulnerability Management (GVM) Architecture and NVT Feed Specification," Greenbone Networks Tech. Rep., 2024.
- [7] S. Zhao and F. Yang, "SIABENCH: Evaluating Agentic AI in Security Incident Analysis and Alert Triage," in *Proc. IEEE Symposium on Security and Privacy (S&P)*, pp. 182-199, 2025.
- [8] J. C. Martinez and R. L. Smith, "Alert Fatigue in Modern Security Operations Centers: A Quantitative Study," *IEEE Transactions on Network and Service Management*, vol. 22, no. 1, pp. 112-125, 2024.
- [9] T. Patel and K. Wilcox, "ForensicLLM: A Fine-Tuned LLaMA-3.1-8B Model for Local Digital Forensic Triage," *IEEE Security and Privacy*, vol. 24, no. 2, pp. 34-42, 2026.
- [10] A. Wu and Y. Zhao, "Multi-Agent Frameworks for Cybersecurity Operations: AutoGen vs. LangGraph," *ACM Transactions on Privacy and Security*, vol. 28, no. 3, pp. 210-225, 2025.
- [11] K. Chen and J. Davis, "Indirect Prompt Injection in Cyber Forensic AI Systems," *IEEE Transactions on Dependable and Secure Computing*, vol. 23, no. 1, pp. 78-92, 2026.
- [12] D. O'Connor and M. Scanlon, "Evaluating AI-Driven Security Incident Response Playbooks," *Journal of Cybersecurity*, vol. 12, no. 1, pp. 55-69, 2025.
- [13] M. Ligh, A. Case, J. Levy, and A. Walters, *The Art of Memory Forensics: Detecting Malware and Threats in Windows, Linux, and Mac Memory*, John Wiley & Sons, 2014.
- [14] R. Johnson and L. Patel, "Evaluating Large Language Models in Digital Investigations: The Problem of Hallucination," *Journal of Digital Forensics, Security and Law*, vol. 19, no. 2, pp. 45-60, 2025.
- [15] National Institute of Standards and Technology (NIST), "Zero Trust Architecture," NIST Special Publication 800-207, 2020.
- [16] P. Mell and T. Grance, "The NIST Definition of Cloud Computing," NIST Special Publication 800-145, 2011.
- [17] S. Schatz, "Plaso: Forensic Timeline Creation and Management," *Digital Investigation*, vol. 20, pp. 12-21, 2017.
- [18] ISO/IEC, "Information technology - Security techniques - Information security management systems - Requirements," ISO/IEC Standard 27001:2022, 2022.
- [19] F. Zhao, "Large Language Models Security and Prompt Injection Mitigation," *IEEE Xplore Digital Library*, 2025.
- [20] K. Patel, "Graph-Based Attack Path Analysis and Lateral Movement Enumeration," *ACM Transactions on Database Systems*, vol. 30, no. 2, pp. 110-124, 2025.
- [21] M. Ramaswamy and E. Al-Shaer, "Automated Threat Modeling and Risk Quantification for Enterprise Networks," *IEEE Transactions on Dependable and Secure Computing*, vol. 21, no. 4, pp. 1820-1834, 2024.
- [22] H. Holm and K. Sommestad, "Using the Common Vulnerability Scoring System to Predict Exploit Likelihood," *IEEE Transactions on Software Engineering*, vol. 42, no. 8, pp. 784-798, 2016.
- [23] S. Nayak and B. Rao, "Explainable AI in Cyber Threat Intelligence: A Comparative Study of LIME and SHAP," *IEEE Transactions on Information Forensics and Security*, vol. 19, pp. 4210-4223, 2024.
- [24] C. Richardson, *Digital Forensic Triage and Incident Response*, Manning Publications, 2024.